


Exploratory Data Analysis (EDA) — Complete & Detailed Cheatsheet

“By Akshen Dhami”

What is EDA?

Exploratory Data Analysis (EDA) is the process of exploring, summarizing, and visualizing data to understand its underlying structure, detect patterns, and find anomalies before applying modeling or prediction.

When to Use:

 Always perform EDA before any modeling or machine learning — it forms your understanding of the dataset and reveals necessary preprocessing steps.

Understanding Data

Data

Units of information — structured (like CSV files or databases) or unstructured (like images, audio, or text).

Types of Variables

1. Numerical Variables (Quantitative)

Type	Description	Example	Python Example
Continuous	Infinite possible values within a range	Temperature, height, weight	tips['total_bill']
Discrete	Countable finite values	Number of siblings, products bought	titanic['sibsp']

```
import seaborn as sns
tips = sns.load_dataset('tips')
titanic = sns.load_dataset('titanic')
print(tips['total_bill'].head()) # Continuous
print(titanic['sibsp'].head()) # Discrete
```

When to Use:

- ✓ Continuous → For statistical and numerical modeling (regression, distribution analysis).
- ✓ Discrete → When counting items, frequencies, or categorical occurrences numerically.

2. Categorical Variables (Qualitative)

Type	Description	Example	Python Code
Nominal	No order among categories	Gender (Male/Female)	titanic['sex'].head()
Binary	Exactly two categories	Smoker: Yes/No	tips['smoker'].head()
Ordinal	Ordered categories	Class: 1st > 2nd > 3rd	titanic['class'].head()

When to Use:

- ✓ Nominal → When categories have **no ranking** (e.g., colors, city names).
- ✓ Ordinal → When categories have **meaningful order** (e.g., education level).
- ✓ Binary → When working with **two possible states** (True/False, 1/0).

Data Types (Check & Change)

Check Data Types

`titanic.dtypes`

Change Data Type

`titanic['age'] = titanic['age'].astype('float')`

When to Use:

- ✓ Always check data types before computation or visualization.
- ✓ Convert incorrect data types (like strings in numeric columns) to appropriate formats to avoid calculation errors.

Accessing Data with loc and iloc

Method	Description	Example	Output
loc	Access rows/columns by labels	<code>titanic.loc[0, 'sex']</code>	Access “sex” of first passenger
iloc	Access by index positions	<code>titanic.iloc[0, 2]</code>	Access first row, third column
loc slice	Select multiple rows by label	<code>titanic.loc[0:5, ['sex', 'age']]</code>	Returns multiple rows and columns
iloc slice	Select by range	<code>titanic.iloc[0:5, 0:3]</code>	First 5 rows, 3 columns

When to Use:

- ✓ `loc` → When you know the **column names** or **row labels**.
- ✓ `iloc` → When you want to slice by **numerical index position**.
- ✓ Always use `iloc` in loops or when column names are dynamic.

Descriptive Statistics

Measure	Description	Example	Python Code
Mean	Arithmetic average	Average restaurant bill	<code>tips['total_bill'].mean()</code>
Median	Middle value	Middle bill value	<code>tips['total_bill'].median()</code>
Mode	Most frequent	Common tip	<code>tips['tip'].mode()[0]</code>

`tips[['total_bill', 'tip']].describe()`

When to Use:

- ✓ Mean → Data is **normally distributed** and without outliers.
- ✓ Median → Data is **skewed** or has **outliers** (robust to extreme values).
- ✓ Mode → Data is **categorical** or you want the most common occurrence.

Distribution of Data

`sns.histplot(tips['total_bill'], kde=True)`

When to Use:

- ✓ Always visualize distribution before deciding on transformations or model types.
- ✓ If highly skewed → apply **log** or **Yeo-Johnson** transformations.

Measures of Dispersion

Measure	Description	Code	When to Use
Variance	Spread of data around mean	<code>tips['total_bill'].var()</code>	To measure variability in continuous data
Standard Deviation	Square root of variance	<code>tips['total_bill'].std()</code>	To compare spread in similar scale
Coefficient of Variation	Std ÷ Mean	<code>tips['total_bill'].std()/tips['total_bill'].mean()</code>	When comparing relative variability across different datasets

Skewness & Kurtosis

Concept	Meaning	Types	Example	Code
Skewness	Asymmetry of data	Right (+ve), Left (-ve), Zero	Income (right-skewed)	<code>tips['total_bill'].skew()</code>
Kurtosis	Peakedness	Leptokurtic (peaked), Platykurtic (flat), Mesokurtic (normal)	Exam scores clustering	<code>tips['total_bill'].kurt()</code>

When to Use:

- ✓ Use skewness & kurtosis to understand **shape of data distribution** before applying parametric statistical tests.
- ✓ If skewness > ±1 → data is **heavily skewed** → consider transformation.

Covariance & Correlation

Measure	Description	Types	Example	Code
Covariance	How two variables vary together	+ve, -ve, 0	Bill ↑ → Tip ↑	<code>tips.cov()</code>
Correlation	Strength of relationship (-1 to +1)	Positive, Negative, None	Bill vs Tip	<code>tips.corr()</code>

`sns.heatmap(tips.corr(), annot=True, cmap='coolwarm')`

When to Use:

- ✓ Correlation → for **linear relationships** between variables.
- ✓ Covariance → for **magnitude and direction** of co-movement.
- ✓ Use **correlation** more often — easier to interpret and scale-independent.

Empirical Rule (68–95–99.7 Rule) & Z-Score

Empirical Rule

- 68% of values within **1 std**
- 95% within **2 std**
- 99.7% within **3 std**

Z-Score

Represents how many std deviations a value is from mean.

from scipy import stats

```
tips['z_score'] = stats.zscore(tips['total_bill'])
```

When to Use:

- ✓ Use Z-scores for **outlier detection** and **standardization**.
- ✓ Great for normally distributed variables to find extremes.

Outliers

What: Data points far from the central trend.

1 Z-Score Method

```
outliers_z = tips[(tips['z_score'] > 3) | (tips['z_score'] < -3)]
```

2 IQR Method

```
Q1, Q3 = tips['tip'].quantile([0.25, 0.75])
```

```
IQR = Q3 - Q1
```

```
outliers_iqr = tips[(tips['tip'] < (Q1 - 1.5 * IQR)) | (tips['tip'] > (Q3 + 1.5 * IQR))]
```

When to Use:

- ✓ Use **Z-score** for **normal data**;
- ✓ Use **IQR** for **non-normal or skewed data**.

Missing Values

```
titanic.isnull().sum()
```

Method	Description	Code	When to Use
Imputation	Replace with mean/median/mode	<code>titanic['age'].fillna(titanic['age'].median(), inplace=True)</code>	If missingness is small and predictable
Deletion	Drop rows/columns	<code>titanic.dropna(subset=['embarked'], inplace=True)</code>	If missing data is small fraction (<5%)
Visualization	Detect pattern	<code>sns.heatmap(titanic.isnull(), cbar=False)</code>	Before imputing, check patterns visually

Encoding Categorical Data

Type	Description	Code	When to Use
Dummy (k-1)	Creates k-1 binary columns	<code>pd.get_dummies(titanic['sex'], drop_first=True)</code>	Use when avoiding multicollinearity

Type	Description	Code	When to Use
One-Hot	Creates separate binary columns	<code>pd.get_dummies(titanic['sex'])</code>	For non-ordinal, independent categories
Label	Converts to integer codes	<code>from sklearn.preprocessing import LabelEncoder; titanic['sex']=LabelEncoder().fit_transform(titanic['sex'])</code>	For ordinal categories or tree models
Frequency	Uses occurrence frequency	<code>freq = titanic['embarked'].value_counts()/len(titanic); titanic['embarked']=titanic['embarked'].map(freq)</code>	When number of categories is large

Feature Scaling

Method	Description	Code	When to Use
Standardization	Mean=0, Std=1	<code>from sklearn.preprocessing import StandardScaler; tips[['total_bill','tip']] = StandardScaler().fit_transform(tips[['total_bill','tip']])</code>	For algorithms like SVM, PCA, Logistic Regression
Min-Max	Scale between 0–1	<code>from sklearn.preprocessing import MinMaxScaler; tips[['total_bill','tip']] = MinMaxScaler().fit_transform(tips[['total_bill','tip']])</code>	For bounded feature importance (Neural Networks, KNN)

Data Transformation

Transformation	Purpose	Clarification	Code	When to Use
Log Transform	Reduce right skewness	Makes data symmetric	<code>import numpy as np; np.log(tips['total_bill'])</code>	For exponential distributions (e.g. income, sales)
Exponential Transform	Inverse of log	Revert to original	<code>np.exp(tips['total_bill'])</code>	To reverse log scaling
Box-Cox	Normalize positive-only data	Requires >0	<code>from scipy import stats; stats.boxcox(tips['total_bill'])</code>	When data is strictly positive

Transformation	Purpose	Clarification	Code	When to Use
Yeo-Johnson	Works with negatives	Handles zeros/negatives	from scipy import stats; stats.yeojohnson(tips['total_bill'])	When data includes negative values

Types of EDA Analysis

Type	Focus	Numerical Methods	Categorical Methods	When to Use
Univariate	1 variable	Histogram, Boxplot, Summary	Count plot, Bar chart	To understand distribution of individual variables
Bivariate	2 variables	Scatterplot, Correlation	Crosstab, Stacked bar	To explore relationships
Multivariate	3+ variables	Pairplot, Heatmap	Grouped bar, Hue plots	To observe combined effects and interactions

Train-Test Split

```
from sklearn.model_selection import train_test_split
```

```
X = tips[['total_bill', 'size']]
```

```
y = tips['tip']
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3)
```

When to Use:

- ✓ Before model training to evaluate generalization.
- ✓ 70–30 or 80–20 split is typical depending on data size.

✓ Summary — Everything Covered

- ✓ Data types & variable types
- ✓ Data type checks & conversions
- ✓ Accessing data (iloc, loc)
- ✓ Descriptive stats
- ✓ Dispersion & shape measures
- ✓ Covariance & correlation
- ✓ Empirical Rule & Z-scores
- ✓ Outlier detection (Z & IQR)
- ✓ Missing value handling
- ✓ Encoding & scaling
- ✓ Data transformation
- ✓ EDA types (univariate/bivariate/multivariate)
- ✓ Train-test split with use cases